

Group Number 1

Members :

- Sunil Jacob
- Sandeep
- Hari Ganapathy
- Sasirekha Sathasivam

Parkinsons Disease Data Set

Dataset information:

This dataset is composed of a range of biomedical voice measurements from 31 people, 23 with Parkinson's disease (PD). Each column in the table is a particular voice measure, and each row corresponds one of 195 voice recording from these individuals ("name" column). The main aim of the data is to discriminate healthy people from those with PD, according to "status" column which is set to 0 for healthy and 1 for PD.

The data is in ASCII CSV format. The rows of the CSV file contain an instance corresponding to one voice recording. There are around six recordings per patient, the name of the patient is identified in the first column.

Attribute Information:

Matrix column entries (attributes): * name - ASCII subject name and recording number * MDVP:Fo(Hz) - Average vocal fundamental frequency * MDVP:Fhi(Hz) - Maximum vocal fundamental frequency * MDVP:Flo(Hz) - Minimum vocal fundamental frequency * MDVP:Jitter(%),MDVP:Jitter(Abs),MDVP:RAP,MDVP:PPQ,Jitter:DDP - Several measures of variation in fundamental frequency * MDVP:Shimmer,MDVP:Shimmer(dB),Shimmer:APQ3,Shimmer:APQ5,MDVP:APQ,Shimmer:DDA - Several measures of variation in amplitude * NHR,HNR - Two measures of ratio of noise to tonal components in the voice * status - Health status of the subject (one) - Parkinson's, (zero) - healthy * RPDE,D2 - Two nonlinear dynamical complexity measures * DFA - Signal fractal scaling exponent * spread1,spread2,PPE - Three nonlinear measures of fundamental frequency variation

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB as gnb
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from IPython.display import Image
from sklearn import tree
from os import system
```

```
data = pd.read_csv('parkinsons.txt')
```

```
data.head()
```

	name	MDVP:Fo(Hz)	MDVP:Fhi(Hz)	MDVP:Flo(Hz)	MDVP:Jitter(%)	MDVP:Jitter(Abs)	MDVP:RAP	MDVP:PPQ	Jitter:DDP	MDVP:Shimmer
0	phon_R01_S01_1	119.992	157.302	74.997	0.00784	0.00007	0.00370	0.00554	0.01109	0.04374
1	phon_R01_S01_2	122.400	148.650	113.819	0.00968	0.00008	0.00465	0.00696	0.01394	0.06134
2	phon_R01_S01_3	116.682	131.111	111.555	0.01050	0.00009	0.00544	0.00781	0.01633	0.05233
3	phon_R01_S01_4	116.676	137.871	111.366	0.00997	0.00009	0.00502	0.00698	0.01505	0.05492
4	phon_R01_S01_5	116.014	141.781	110.655	0.01284	0.00011	0.00655	0.00908	0.01966	0.06425

5 rows x 24 columns

```
data.shape
```

```
(195, 24)
```

```
#moving status column to the last
df1=data.pop('status')
data['status'] = df1
```

```
data.describe().transpose()
```


	count	mean	std	min	25%	50%	75%	max
MDVP:Fo(Hz)	195.0	154.228641	41.390065	88.333000	117.572000	148.790000	182.769000	260.105000
MDVP:Fhi(Hz)	195.0	197.104918	91.491548	102.145000	134.862500	175.829000	224.205500	592.030000
MDVP:Flo(Hz)	195.0	116.324631	43.521413	65.476000	84.291000	104.315000	140.018500	239.170000
MDVP:Jitter(%)	195.0	0.006220	0.004848	0.001680	0.003460	0.004940	0.007365	0.033160
MDVP:Jitter(Abs)	195.0	0.000044	0.000035	0.000007	0.000020	0.000030	0.000060	0.000260
MDVP:RAP	195.0	0.003306	0.002968	0.000680	0.001660	0.002500	0.003835	0.021440
MDVP:PPQ	195.0	0.003446	0.002759	0.000920	0.001860	0.002690	0.003955	0.019580
Jitter:DDP	195.0	0.009920	0.008903	0.002040	0.004985	0.007490	0.011505	0.064330
MDVP:Shimmer	195.0	0.029709	0.018857	0.009540	0.016505	0.022970	0.037885	0.119080
MDVP:Shimmer(dB)	195.0	0.282251	0.194877	0.085000	0.148500	0.221000	0.350000	1.302000
Shimmer:APQ3	195.0	0.015664	0.010153	0.004550	0.008245	0.012790	0.020265	0.056470
Shimmer:APQ5	195.0	0.017878	0.012024	0.005700	0.009580	0.013470	0.022380	0.079400
MDVP:APQ	195.0	0.024081	0.016947	0.007190	0.013080	0.018260	0.029400	0.137780
Shimmer:DDA	195.0	0.046993	0.030459	0.013640	0.024735	0.038360	0.060795	0.169420
NHR	195.0	0.024847	0.040418	0.000650	0.005925	0.011660	0.025640	0.314820
HNR	195.0	21.885974	4.425764	8.441000	19.198000	22.085000	25.075500	33.047000
RPDE	195.0	0.498536	0.103942	0.256570	0.421306	0.495954	0.587562	0.685151
DFA	195.0	0.718099	0.055336	0.574282	0.674758	0.722254	0.761881	0.825288
spread1	195.0	-5.684397	1.090208	-7.964984	-6.450096	-5.720868	-5.046192	-2.434031
spread2	195.0	0.226510	0.083406	0.006274	0.174351	0.218885	0.279234	0.450493
D2	195.0	2.381826	0.382799	1.423287	2.099125	2.361532	2.636456	3.671155
PPE	195.0	0.206552	0.090119	0.044539	0.137451	0.194052	0.252980	0.527367
status	195.0	0.753846	0.431878	0.000000	1.000000	1.000000	1.000000	1.000000

```
# There are 147 people affected with PD .
data.status.value_counts()
```

```
1    147
0     48
Name: status, dtype: int64
```

```
sns.pairplot(data)
```


```
<seaborn.axisgrid.PairGrid at 0x117064278>
```



Univariate Analysis

```
fig, ax = plt.subplots(1,3,figsize=(16,10))
sns.boxplot(x='spread1',data=data, ax=ax[0],orient='v')
sns.boxplot(x='spread2',data=data, ax=ax[1],orient='v')
sns.boxplot(x='PPE',data=data, ax=ax[2],orient='v')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x119816f60>
```



The above figure shows the box plot of the frequency variation. All the three variations have outliers.

Generally speaking, decision trees are able to handle outliers. It is very unlikely that decision tree will create a leaf to isolate them

```
# dropping name column as this column is not much significant
data = data.drop('name',axis=1)
```

```
type(data)
```

```
pandas.core.frame.DataFrame
```


```
# Using z-score function in scipy to detect outliers
# from scipy import stats
# z = np.abs(stats.zscore(data))
# print(z)
```

```
# Defining a threshold to detect an outlier as we cannot infer any meaningful info from the above dataset
# threshold = 3
# print(np.where(z > 3))
```

```
# print(z[73][1])
# data = data[(z < 3).all(axis=1)]
```

```
fig, ax = plt.subplots(1,3,figsize=(16,8))
sns.distplot(data['MDVP:Flo(Hz)'],ax=ax[0])
sns.distplot(data['MDVP:Fo(Hz)'],ax=ax[1])
sns.distplot(data['MDVP:Phi(Hz)'],ax=ax[2])
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x119d37dd8>
```



The measures of vocal fundamental frequency are shown above. There is a positive skewness for `minimum vocal fundamental frequency` with more high values between 75Hz and 125Hz. The `average vocal frequency` is almost normally distributed with more values ranging 115Hz and 130Hz. The `high vocal frequency` does not have any skewness, but some range of values are at the right most tail

```
fig, ax = plt.subplots(1,2,figsize=(16,8))
sns.distplot(data['NHR'],ax=ax[0])
sns.distplot(data['HNR'],ax=ax[1])
```

<matplotlib.axes._subplots.AxesSubplot at 0x11a3fe908>




The measure of tonal component of frequency is shown above. The value `NHR` is right skewed for there are so many observations in the area, but they seem to be with very minimal values. The maximum number of observations is between 0 and 0.04.

The value `HNR` looks like normally distributed, but in a first look there seems to be a slight negative skewness

```
fig, ax = plt.subplots(2,3,figsize=(16,8))
sns.distplot(data['MDVP:Shimmer'],ax=ax[0,0])
sns.distplot(data['MDVP:Shimmer(dB)'],ax=ax[0,1])
sns.distplot(data['Shimmer:APQ3'],ax=ax[0,2])
sns.distplot(data['Shimmer:APQ5'],ax=ax[1,0])
sns.distplot(data['MDVP:APQ'],ax=ax[1,1])
sns.distplot(data['Shimmer:DDA'],ax=ax[1,2])
```

<matplotlib.axes._subplots.AxesSubplot at 0x11a9a4be0>




For all of the above graphs, we can observe that the measure of variation in amplitude is positively skewed

Bi -Variate Analysis

Relation between target variable and independent variables

```
sns.distplot( data[data.status == 0]['spread1'], color = 'r')
sns.distplot( data[data.status == 1]['spread1'], color = 'g')
```


<matplotlib.axes._subplots.AxesSubplot at 0x11a4f07f0>



spread1 is normally distributed between person who have PD and who is normal. People who have **spread1** between - 8.5 and -7.5 are more and they are normal. People whose **spread1** is between -6.5 and -5 are having PD

```
sns.distplot( data[data.status == 0]['spread2'], color = 'r')
sns.distplot( data[data.status == 1]['spread2'], color = 'g')
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x11af8fe48>
```



spread1 and **spread2** is normally distributed. between person who have PD and who is normal. People who have **spread1** between 0.15 and 0.175 are more and they are normal. People whose **spread1** is between 0.175 and 0.2 are having PD



```
sns.distplot( data[data.status == 0]['PPE'], color = 'r')
sns.distplot( data[data.status == 1]['PPE'], color = 'g')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x11b0d7e80>
```



```
fig, ax = plt.subplots(1,2,figsize=(16,8))
sns.boxplot(x='status',y='NHR',data=data,ax=ax[0])
sns.boxplot(x='status',y='HNR',data=data,ax=ax[1])
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x11b403390>
```



People who have PD(status equal to one) have higher levels of Noise to Harmonic ratio. Also, looking into the HNR ratio people who have PD have lower levels in the same.


```
sns.boxplot(x='status',y='MDVP:Phi(Hz)',data=data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x11b5115c0>
```




```
fig, ax = plt.subplots(1,2,figsize=(16,8))
sns.boxplot(x='status',y='MDVP:Flo(Hz)',data=data,palette="Set1",ax=ax[0])
sns.boxplot(x='status',y='MDVP:F0(Hz)',data=data,palette="Set1",ax=ax[1])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x11b88b0b8>
```

When we look the relationship between status and MDVP:F0(Hz) we can see the median value is around 199 Hz for people who are normal. For people who are affected with Parkinsons the median value comes around 145 Hz


```
# For categorical predictors
cols = ["MDVP:Jitter(%)", "MDVP:Jitter(%)", "MDVP:RAP", "MDVP:PPQ", "Jitter:DDP"]
fig, axs = plt.subplots(ncols = 5, figsize=(16,8))
fig.tight_layout()
for i in range(0,len(cols)):
    sns.boxplot(x='status',y=cols[i],data=data, ax = axs[i])
```



People who are suffering for PD tend to have higher jitter %. It seems if the values goes above 0.15 we can confirm the patient is having PD. The variation of fundamental frequency is in a low range for people who is normal.

Correlation comparison with heat map

```
import matplotlib.pyplot as plt
import seaborn as sns
corr = data.corr()
sns.set_context("notebook", font_scale=1.0, rc={"lines.linewidth": 3.5})
plt.figure(figsize=(18,7))
# create a mask so we only see the correlation values once
mask = np.zeros_like(corr)
mask[np.triu_indices_from(mask, 1)] = True
a = sns.heatmap(corr,mask=mask, annot=True, fmt='.2f')
rotx = a.set_xticklabels(a.get_xticklabels(), rotation=90)
roty = a.set_yticklabels(a.get_yticklabels(), rotation=30)
```



- **MDVP:Jitter(%)** has a very high correlation with **MDVP:Jitter(Abs)**,**MDVP:RAP**,**MDVP:PPQ**,**Jitter:DDP**
- **MDVP:Shimmer** has a very correlation with **MDVP:Shimmer(dB)**,**Shimmer:APQ3**,**Shimmer:APQ5**,**MDVP:APQ**,**Shimmer:DDA** this may be because they are related to each other. This may be because `multi-dimensinal voice programs analysis` is closely related with these variables
- The target variable **status** has a weak positive correlation with **spread1**

Applying models

Decision Tree

Decision trees can be used to predict both continuous and discrete values i.e. they work well for both regression and classification tasks.

```
from sklearn import metrics
```

```
X = data.drop("status",axis=1)
Y = data["status"]
```

```
# Splitting Data into 70% Training data and 30% Testing Data:
X_train, X_test, y_train, y_test = train_test_split(X, Y,train_size=0.7, test_size=0.3, random_state=42)
print(len(X_train)),print(len(X_test))
```

```
136
59
```

```
(None, None)
```

```
# Applying decision tree model
dt_model = DecisionTreeClassifier(criterion='entropy',max_depth=6,random_state=100,min_samples_leaf=5)
```

```
dt_model.fit(X_train, y_train)
```

```
DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=6,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_split=1e-07, min_samples_leaf=5,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        presort=False, random_state=100, splitter='best')
```

```
dt_model.score(X_test , y_test)
```

```
0.89830508474576276
```

```
y_pred = dt_model.predict(X_test)
```

```
confusion_matrix(y_test,y_pred)
```

```
array([[11,  4],
       [ 2, 42]])
```

```
#Count mis-classified one
count_misclassified = (y_test != y_pred).sum()
print('Misclassified samples: {}'.format(count_misclassified))
```

```
Misclassified samples: 6
```

```
from IPython.display import Image
from sklearn import tree
from os import system

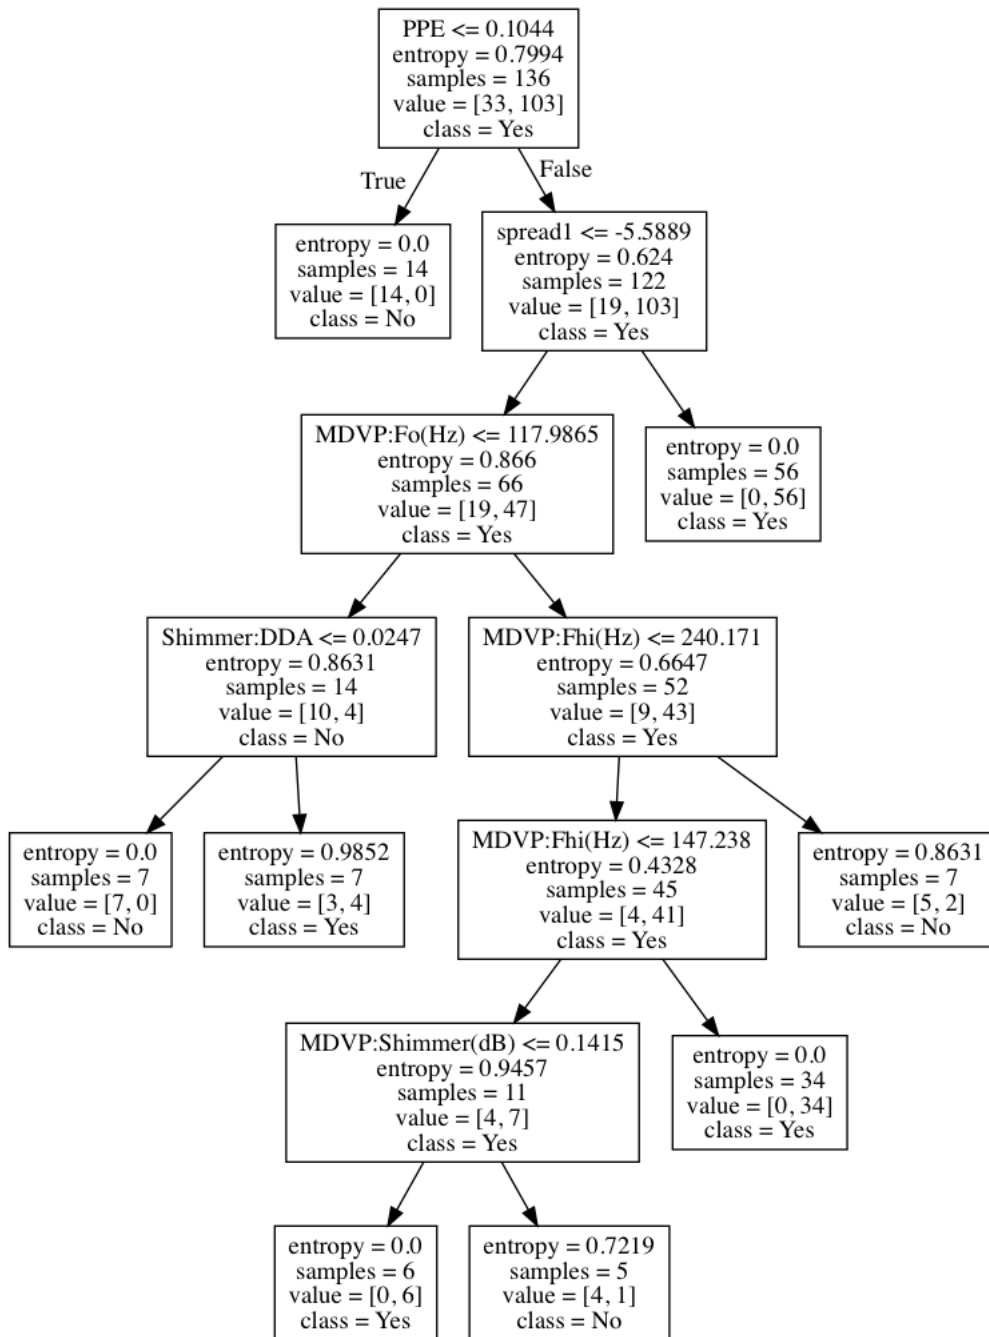
train_char_label = ['No', 'Yes']
pd_tree_regularized = open('pd_tree_regularized.dot','w')
dot_data = tree.export_graphviz(dt_model, out_file= pd_tree_regularized , feature_names = list(X_train), class_names = list(train_char_label))

pd_tree_regularized.close()

print (pd.DataFrame(dt_model.feature_importances_ , columns = ["Imp"], index = X_train.columns))
```

	Imp
MDVP:Fo(Hz)	0.113981
MDVP:Fhi(Hz)	0.196589
MDVP:Flo(Hz)	0.000000
MDVP:Jitter(%)	0.000000
MDVP:Jitter(Abs)	0.000000
MDVP:RAP	0.000000
MDVP:PPQ	0.000000
Jitter:DDP	0.000000
MDVP:Shimmer	0.000000
MDVP:Shimmer(dB)	0.073695
Shimmer:APQ3	0.000000
Shimmer:APQ5	0.000000
MDVP:APQ	0.000000
Shimmer:DDA	0.056276
NHR	0.000000
HNR	0.000000
RPDE	0.000000
DFA	0.000000
spread1	0.205876
spread2	0.000000
D2	0.000000
PPE	0.353584

```
# You can also copy the script in the .dot file and paste it at http://webgraphviz.com/ to get tree view
# or create a .png as below
system("dot -Tpng pd_tree_regularized.dot -o pd_tree_regularized.png")
Image("pd_tree_regularized.png")
```



K Nearest Neighbour

```
k_model = KNeighborsClassifier(n_neighbors=5)
k_model.fit(X_train, y_train)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=1, n_neighbors=5, p=2,
                     weights='uniform')
```

```
k_model.score(X_test, y_test)
```

```
0.83050847457627119
```

```
y_pred = k_model.predict(X_test)
```

```
count_misclassified = (y_test != y_pred).sum()
print('Misclassified samples in KNN: {}'.format(count_misclassified))
```

```
Misclassified samples in KNN: 10
```

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
rfcl = RandomForestClassifier(n_estimators = 50)
rfcl = rfcl.fit(X_train, y_train)
```

```
y_pred = rfcl.predict(X_test)
rfcl.score(X_test , y_test)
```


```
0.94915254237288138
```

```
count_misclassified = (y_test != y_pred).sum()
print('Misclassified samples in Random Forest: {}'.format(count_misclassified))
```

```
Misclassified samples in Random Forest: 3
```

```
feature_imp = pd.Series(rfcl.feature_importances_,index=X.columns).sort_values(ascending=False)
feature_imp
# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

```
/Users/suniljacob/anaconda/lib/python3.6/site-packages/matplotlib/axes/_axes.py:545: UserWarning: No labelled objects found. Use label='...' kwa
rg on individual plots.
  warnings.warn("No labelled objects found. ")
```



Bagging

```
from sklearn.ensemble import BaggingClassifier
bgcl = BaggingClassifier(base_estimator=dt_model, n_estimators=50, max_samples=.7)
bgcl = bgcl.fit(X_train, y_train)
y_pred = bgcl.predict(X_test)
bgcl.score(X_test , y_test)
```

0.89830508474576276

```
count_misclassified = (y_test != y_pred).sum()
print('Misclassified samples in Bagging: {}'.format(count_misclassified))
```

Misclassified samples in Bagging: 6

AdaBoosting

```
from sklearn.ensemble import AdaBoostClassifier
abcl = AdaBoostClassifier( n_estimators= 50)
abcl = abcl.fit(X_train,y_train)
y_pred = abcl.predict(X_test)
abcl.score(X_test , y_test)
```

0.89830508474576276

```
count_misclassified = (y_test != y_pred).sum()
print('Misclassified samples in Ada Boosting: {}'.format(count_misclassified))
```

Misclassified samples in Ada Boosting: 6

Gradient Boosting

```
from sklearn.ensemble import GradientBoostingClassifier
gbcl = GradientBoostingClassifier(n_estimators = 50, learning_rate = 0.05)
gbcl = gbcl.fit(X_train,y_train)
y_pred = gbcl.predict(X_test)
gbcl.score(X_test , y_test)
```

0.89830508474576276

```
count_misclassified = (y_test != y_pred).sum()
print('Misclassified samples in Gradient Boosting: {}'.format(count_misclassified))
```

Misclassified samples in Gradient Boosting: 6

Of all the above ones Random Forest algorithm gave the maximum accuracy.